



Plant Pathology (Image Classification)



Plant_Pathology Dataset

- Given a photo of an apple leaf, can you accurately assess its health? This competition will challenge you to distinguish between leaves which are healthy, those which are infected with apple rust, those that have apple scab, and those with more than one disease.
- A folder containing the train and test images, in jpg format.



1-PreProcessing

Read Data

```
In [ ]: train_df = pd.read_csv("/kaggle/input/plant-pathology-2020-fgvc7/train.csv")
train_df['img_name'] = train_df['image_id'] + ".jpg"
train_df.head()
```

Out[4]:

	image_id	healthy	multiple_diseases	rust	scab	img_name
0	Train_0	0	0	0	1	Train_0.jpg
1	Train_1	0	1	0	0	Train_1.jpg
2	Train_2	1	0	0	0	Train_2.jpg
3	Train_3	0	0	1	0	Train_3.jpg
4	Train_4	1	0	0	0	Train_4.jpg

Determine Directions

```
In [ ]: from shutil import copyfile

# delete temp dir
if os.path.exists('/kaggle/temp/'):
    shutil.rmtree('/kaggle/temp/')

os.mkdir('/kaggle/temp/')

# train directory
os.mkdir('/kaggle/temp/train')
os.mkdir('/kaggle/temp/train/healthy')
os.mkdir('/kaggle/temp/train/multiple_diseases')
os.mkdir('/kaggle/temp/train/rust')
os.mkdir('/kaggle/temp/train/scab')

# validation directory
os.mkdir('/kaggle/temp/valid')
os.mkdir('/kaggle/temp/valid/healthy')
os.mkdir('/kaggle/temp/valid/multiple_diseases')
os.mkdir('/kaggle/temp/valid/rust')
os.mkdir('/kaggle/temp/valid/scab')
```

```
In [ ]: SOURCE = '/kaggle/input/plant-pathology-2020-fgvc7/images/'

TRAIN_DIR = '/kaggle/temp/train/'

# copy images to train directory
for index, data in train_set.iterrows():
    label = df.columns[np.argmax(data)]
    filepath = os.path.join(SOURCE, index + ".jpg")
    destination = os.path.join(TRAIN_DIR, label, index + ".jpg")
    copyfile(filepath, destination)

for subdir in os.listdir(TRAIN_DIR):
    print(subdir, len(os.listdir(os.path.join(TRAIN_DIR, subdir))))

healthy 416
scab 465
multiple_diseases 73
rust 502
```

```
In [ ]: VALID_DIR = '/kaggle/temp/valid/'

# copy images to valid directory
for index, data in valid_set.iterrows():
    label = df.columns[np.argmax(data)]
    filepath = os.path.join(SOURCE, index + ".jpg")
    destination = os.path.join(VALID_DIR, label, index + ".jpg")
    copyfile(filepath, destination)

for subdir in os.listdir(VALID_DIR):
    print(subdir, len(os.listdir(os.path.join(VALID_DIR, subdir))))

healthy 100
scab 127
multiple_diseases 18
rust 120
```

```
In [ ]: healthy_dir = os.path.join(TRAIN_DIR, 'healthy')
mdiseases_dir = os.path.join(TRAIN_DIR, 'multiple_diseases')
scab_dir = os.path.join(TRAIN_DIR, 'scab')
rust_dir = os.path.join(TRAIN_DIR, 'rust')

healthy_files = os.listdir(healthy_dir)
mdiseases_files = os.listdir(mdiseases_dir)
scab_files = os.listdir(scab_dir)
rust_files = os.listdir(rust_dir)
```

Data Visualization

```
In [ ]: %matplotlib inline

import matplotlib.pyplot as plt
import matplotlib.image as mpimg

pic_index = 2

next_healthy = [os.path.join(healthy_dir, fname) for fname in healthy_files[pic_index-2:pic_index]]
next_mdiseases = [os.path.join(mdiseases_dir, fname) for fname in mdiseases_files[pic_index-2:pic_index]]
next_scab = [os.path.join(scab_dir, fname) for fname in scab_files[pic_index-2:pic_index]]
next_rust = [os.path.join(rust_dir, fname) for fname in rust_files[pic_index-2:pic_index]]

nrows = 4
ncols = 4


fig = plt.gcf()
fig.set_size_inches(ncols*4, nrows*4)

for i, img_path in enumerate(next_healthy+next_mdiseases+next_scab+next_rust):
    sp = plt.subplot(nrows, ncols, i + 1)
    sp.axis('Off') # Don't show axes (or gridlines)
    img = mpimg.imread(img_path)
    plt.title(img_path.split('/')[-2])
    plt.imshow(img)

plt.show()
```



```
In [ ]: def show_imgs(df, num):
    fig, ax = plt.subplots(1,num, figsize=(18,9))
    for x, y in df:
        for img in range(num):
            ax[img].imshow(x[img])
            if y[img][0]:
                title="healthy"
            elif y[img][1]:
                title='multiple_diseases'
            elif y[img][2]:
                title='rust'
            elif y[img][3]:
                title='scab'
            ax[img].set_title(title)
        break
```



Data Balance (Over Sampling)

```
: target_multi_cols = ['healthy', 'multiple_diseases', 'rust', 'scab']

print("Multi Classification Targets")
print(train_df[target_multi_cols].sum())
```

```
Multi Classification Targets
healthy          516
multiple_diseases  91
rust             622
scab             592
dtype: int64
```

Maka data balance (over sampling)

```
: def balance_set(df, x_cols, y_cols):
    ros = RandomOverSampler(random_state=42)

    x_multi, y_multi = ros.fit_resample(df[x_cols], df[y_cols].values)
    data = pd.concat([x_multi, pd.DataFrame(y_multi, columns= y_cols)], axis=1)
    return data

train_multi = balance_set(train_df,
                           x_cols = ["image_id", "img_name"],
                           y_cols = target_multi_cols)

labels = train_multi[target_multi_cols]
label_names = labels[labels==1].stack().reset_index()['level_1']
label_names.index = train_multi.index
train_multi['label_names'] = label_names

print("Multi Classification Labels")
print(train_multi[target_multi_cols].sum())
```

```
Multi Classification Labels
healthy          622
multiple_diseases  622
rust             622
scab             622
dtype: int64
```



Data Split

(Train&Validation)

```
from sklearn.model_selection import train_test_split

train_set, valid_set = train_test_split(df, test_size=0.2, random_state=42)

print(train_set.shape)
print(valid_set.shape)

(1456, 4)
(365, 4)
```

```
tf.random.set_seed(99)
img_data_generator = ImageDataGenerator(rescale=1/255,
                                       validation_split=0.2,
                                       rotation_range = 180,
                                       horizontal_flip = True,
                                       vertical_flip = True,
                                       preprocessing_function=blur_preprocessing
                                       )
```


Image Generator

Before Balance

```
import tensorflow as tf
from tensorflow import keras
import tensorflow_datasets as tfds
import os
import cv2
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
import tensorflow.keras.layers as tfl
import pandas as pd
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras.layers.experimental.preprocessing import RandomFlip, RandomRotation
from PIL import Image
import csv
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

training_datagen = ImageDataGenerator(rescale = 1./255,
                                      rotation_range=40,
                                      width_shift_range=0.2,
                                      height_shift_range=0.2,
                                      shear_range=0.2,
                                      zoom_range=0.2,
                                      horizontal_flip=True,
                                      fill_mode='nearest')

validation_datagen = ImageDataGenerator(rescale = 1./255)

#test_datagen = ImageDataGenerator( rescale = 1.0/255. )

train_generator = training_datagen.flow_from_directory(TRAIN_DIR, target_size=(150,150), class_mode='categorical', batch_size=32)
validation_generator = validation_datagen.flow_from_directory(VALID_DIR, target_size=(150,150), class_mode='categorical', batch_
```

Found 1456 images belonging to 4 classes.
Found 365 images belonging to 4 classes.

After Balance

```
def blur_preprocessing(img):
    return cv2.blur(img, (5, 5))

tf.random.set_seed(99)
img_data_generator = ImageDataGenerator(rescale=1/255,
                                       validation_split=0.2,
                                       rotation_range = 180,
                                       horizontal_flip = True,
                                       vertical_flip = True,
                                       preprocessing_function=blur_preprocessing
                                       )

train_data_multi = img_data_generator.flow_from_dataframe(dataframe=train_multi,
                                                         directory="/kaggle/input/plant-pathology-2020-fgvc7/images/",
                                                         x_col="img_name",
                                                         y_col="label_names",
                                                         target_size=(256, 256),
                                                         class_mode='categorical',
                                                         batch_size=32,
                                                         subset='training',
                                                         shuffle=True,
                                                         seed=42)

val_data_multi = img_data_generator.flow_from_dataframe(dataframe=train_multi,
                                                         directory="/kaggle/input/plant-pathology-2020-fgvc7/images/",
                                                         x_col="img_name",
                                                         y_col="label_names",
                                                         target_size=(256, 256),
                                                         class_mode='categorical',
                                                         batch_size=32,
                                                         subset='validation',
                                                         shuffle=True,
                                                         seed=42)
```

Found 1991 validated image filenames belonging to 4 classes.
Found 497 validated image filenames belonging to 4 classes.



```
def blur_preprocessing(img):  
    return cv2.blur(img, (5, 5))
```

Learning rate scheduler

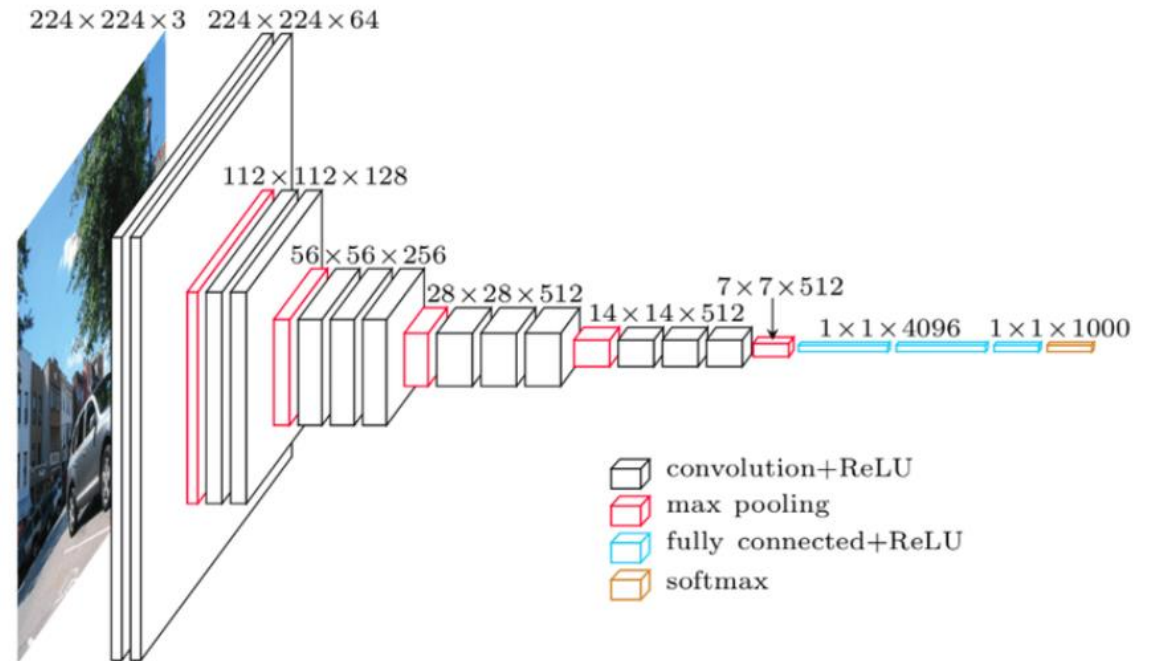
[illegible]

2-Image Classification Models

1- VGG16 (Pretrained Model)

- What Is VGG16?
 - VGG16 refers to the VGG model, also called VGGNet. It is a convolution neural network (CNN) model supporting 16 layers.
 - The VGG16 model can achieve a test accuracy of 92.7% in ImageNet, a dataset containing more than 14 million training images across 1000 object classes. It is one of the top models from the ILSVRC-2014 competition.

VGG16 Architecture



VGG16 modeling

(Transfer learning – compile – fit – load model – evaluation)

```
from tensorflow.keras.applications.vgg16 import VGG16
base_model = VGG16(input_shape = (150, 150, 3), # Shape of our images
include_top = False, # Leave out the last fully connected layer
weights = 'imagenet')
```

```
import tensorflow as tf
import keras
from keras import layers
for layer in base_model.layers:
    layer.trainable = False
    tf.keras.layers.BatchNormalization()
```

```
# Flatten the output layer to 1 dimension
x = layers.Flatten()(base_model.output)
# Add a fully connected layer with 512 hidden units and ReLU activation
x = layers.Dense(512, activation='relu')(x)
# Add a dropout rate of 0.5
x = layers.Dropout(0.5)(x)
# Add a final sigmoid layer with 1 node for classification output
x = layers.Dense(4, activation='softmax')(x)
model = tf.keras.models.Model(base_model.input, x)
model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics = ['acc'])
early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=5)
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint("vgg16.h5", save_best_only=True)
```

```
vgghist = model.fit(train_generator, validation_data = validation_generator, steps_per_epoch = 10, epochs = 15, validation_s
```

Visualize train and validation accuracy

```
import matplotlib.pyplot as plt
acc = vgghist.history['acc']
val_acc = vgghist.history['val_acc']
loss = vgghist.history['loss']
val_loss = vgghist.history['val_loss']
epochs = range(len(acc))
plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend(loc=0)
plt.show()
```

Evaluation

```
model = tf.keras.models.load_model("vgg16.h5") # rollback to best model
model.evaluate(validation_generator)
```

VGG16 modeling

(read test data – make predictions)

```
In [ ]: test_set = pd.read_csv("/kaggle/input/plant-pathology-2020-fgvc7/test.csv", index_col=0)

X_test = []
for index, data in test_set.iterrows():
    filepath = os.path.join(SOURCE, index + ".jpg")
    img = image.load_img(filepath, target_size=(150, 150))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    X_test.append(x)

X_test = np.vstack(X_test) / 255 # rescale images
```

predict the test data

```
In [ ]: y_pred = model.predict(X_test, batch_size=10)
df_out = pd.concat([test_set.reset_index(), pd.DataFrame(y_pred, columns = train_generator.class_indices.keys())], axis=1).set_index('image_id')
df_out.to_csv('submission.csv')
df_out.head()
```

183/183 [=====] - 7s 30ms/step

Out[19]:

	healthy	multiple_diseases	rust	scab
image_id				
Test_0	0.106944	0.044960	0.239254	0.608842
Test_1	0.199255	0.047042	0.460560	0.293144
Test_2	0.014253	0.024624	0.089598	0.871525
Test_3	0.552810	0.033190	0.152912	0.261087
Test_4	0.187862	0.061119	0.353388	0.397631

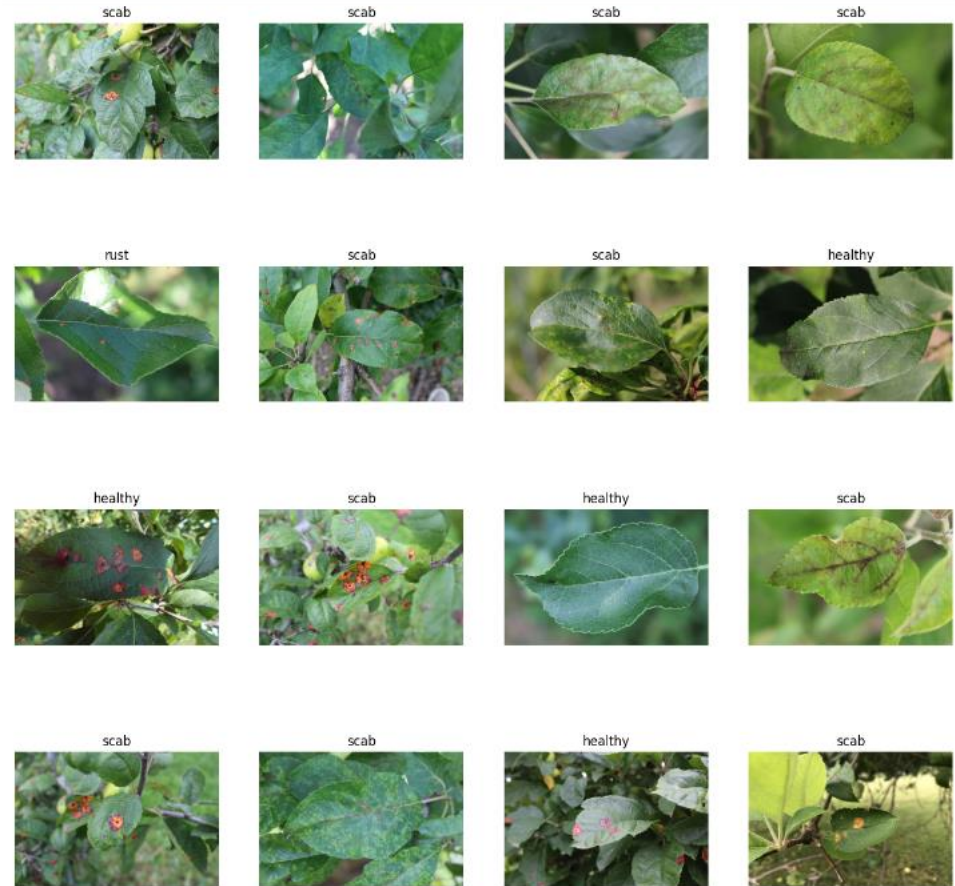
VGG16 modeling

(Predictions)

```
ncols = 4

fig = plt.gcf()
fig.set_size_inches(ncols*4, nrows*4)

for i, (idx, row) in enumerate(df_out.sample(nrows*ncols).iterrows()):
    filepath = os.path.join(SOURCE, idx + ".jpg")
    sp = plt.subplot(nrows, ncols, i + 1)
    sp.axis('Off') # Don't show axes (or gridlines)
    img = mpimg.imread(filepath)
    plt.title(df_out.columns[np.argmax(row)])
    plt.imshow(img)
```



2- Second Model (Build the model)

Second model (Created model)

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(64, (3,3), activation='relu', input_shape=(150, 150, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(128, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(128, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(4, activation='softmax')
])
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 148, 148, 64)	1792
max_pooling2d_4 (MaxPooling 2D)	(None, 74, 74, 64)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	36928
max_pooling2d_5 (MaxPooling 2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_6 (MaxPooling 2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_7 (MaxPooling 2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dropout_1 (Dropout)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3211776
dense_3 (Dense)	(None, 4)	2052
=====		
Total params: 3,473,988		
Trainable params: 3,473,988		
Non-trainable params: 0		

Second Model

(compile - fit)

```
model.compile(loss = 'categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=5)
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint("cm.h5", save_best_only=True)

history = model.fit(train_generator, epochs=50, steps_per_epoch=46,
                    validation_data = validation_generator, validation_steps=12, callbacks=[early_stopping_cb, checkpoint_cb])
```

```
Epoch 1/50
46/46 [=====] - 29s 594ms/step - loss: 0.6115 - accuracy: 0.7782 - val_loss: 0.4194 - val_accuracy: 0.8384
Epoch 2/50
46/46 [=====] - 27s 589ms/step - loss: 0.5390 - accuracy: 0.8043 - val_loss: 0.3652 - val_accuracy: 0.8712
Epoch 3/50
46/46 [=====] - 28s 606ms/step - loss: 0.5331 - accuracy: 0.8015 - val_loss: 0.5373 - val_accuracy: 0.7808
Epoch 4/50
46/46 [=====] - 28s 606ms/step - loss: 0.5130 - accuracy: 0.8084 - val_loss: 0.5944 - val_accuracy: 0.7479
Epoch 5/50
46/46 [=====] - 27s 597ms/step - loss: 0.4785 - accuracy: 0.8310 - val_loss: 0.3099 - val_accuracy: 0.8877
Epoch 6/50
46/46 [=====] - 28s 609ms/step - loss: 0.5099 - accuracy: 0.8159 - val_loss: 0.5052 - val_accuracy: 0.7863
Epoch 7/50
46/46 [=====] - 28s 611ms/step - loss: 0.5072 - accuracy: 0.8187 - val_loss: 0.3551 - val_accuracy: 0.8822
Epoch 8/50
46/46 [=====] - 28s 609ms/step - loss: 0.4575 - accuracy: 0.8400 - val_loss: 0.4851 - val_accuracy: 0.8192
Epoch 9/50
46/46 [=====] - 27s 588ms/step - loss: 0.4559 - accuracy: 0.8455 - val_loss: 0.4275 - val_accuracy: 0.8548
Epoch 10/50
46/46 [=====] - 27s 583ms/step - loss: 0.4740 - accuracy: 0.8468 - val_loss: 0.2493 - val_accuracy: 0.9288
Epoch 11/50
46/46 [=====] - 27s 584ms/step - loss: 0.4095 - accuracy: 0.8571 - val_loss: 0.3297 - val_accuracy: 0.8795
Epoch 12/50
46/46 [=====] - 27s 592ms/step - loss: 0.4077 - accuracy: 0.8516 - val_loss: 0.2552 - val_accuracy: 0.9151
Epoch 13/50
46/46 [=====] - 28s 618ms/step - loss: 0.3895 - accuracy: 0.8723 - val_loss: 0.2987 - val_accuracy: 0.8904
Epoch 14/50
46/46 [=====] - 27s 591ms/step - loss: 0.4016 - accuracy: 0.8654 - val_loss: 0.4763 - val_accuracy: 0.8438
Epoch 15/50
46/46 [=====] - 28s 603ms/step - loss: 0.3788 - accuracy: 0.8764 - val_loss: 0.2368 - val_accuracy: 0.9260
Epoch 16/50
46/46 [=====] - 28s 602ms/step - loss: 0.3448 - accuracy: 0.8832 - val_loss: 0.3211 - val_accuracy: 0.9041
Epoch 17/50
46/46 [=====] - 26s 575ms/step - loss: 0.3652 - accuracy: 0.8750 - val_loss: 0.2832 - val_accuracy: 0.8986
Epoch 18/50
46/46 [=====] - 27s 587ms/step - loss: 0.3788 - accuracy: 0.8757 - val_loss: 0.2590 - val_accuracy: 0.9178
Epoch 19/50
46/46 [=====] - 27s 582ms/step - loss: 0.3467 - accuracy: 0.8853 - val_loss: 0.3358 - val_accuracy: 0.8822
Epoch 20/50
46/46 [=====] - 27s 595ms/step - loss: 0.3446 - accuracy: 0.8839 - val_loss: 0.2467 - val_accuracy: 0.9151
```


Second Model

(train and validation accuracy visualization – load model – evaluation)

```
In [ ]: import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend(loc=0)
```



```
In [ ]: model = tf.keras.models.load_model("cm.h5") # rollback to best model
model.evaluate(validation_generator)
```

12/12 [=====] - 4s 337ms/step - loss: 0.2368 - accuracy: 0.9260

Out[25]: [0.2368302196264267, 0.9260274171829224]

Second Model

(read test data – make predictions)

```
In [ ]: test_set = pd.read_csv("/kaggle/input/plant-pathology-2020-fgvc7/test.csv", index_col=0)

X_test = []
for index, data in test_set.iterrows():
    filepath = os.path.join(SOURCE, index + ".jpg")
    img = image.load_img(filepath, target_size=(150, 150))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    X_test.append(x)

X_test = np.vstack(X_test) / 255 # rescale images
```

```
In [ ]: y_pred = model.predict(X_test, batch_size=10)
df_out = pd.concat([test_set.reset_index(), pd.DataFrame(y_pred, columns = train_generator.class_indices.keys()), axis=1).set_index('image_id')
df_out.to_csv('submission.csv')
df_out.head()
```

183/183 [=====] - 1s 4ms/step

Out[27]:

	healthy	multiple_diseases	rust	scab
image_id				
Test_0	6.722297e-12	0.000572	0.999428	7.045307e-13
Test_1	1.999887e-08	0.015762	0.984238	1.652383e-08
Test_2	4.754157e-02	0.017981	0.000729	9.337487e-01
Test_3	8.150271e-01	0.023522	0.002990	1.584612e-01
Test_4	1.529118e-18	0.001206	0.998794	3.132368e-18

Second Model

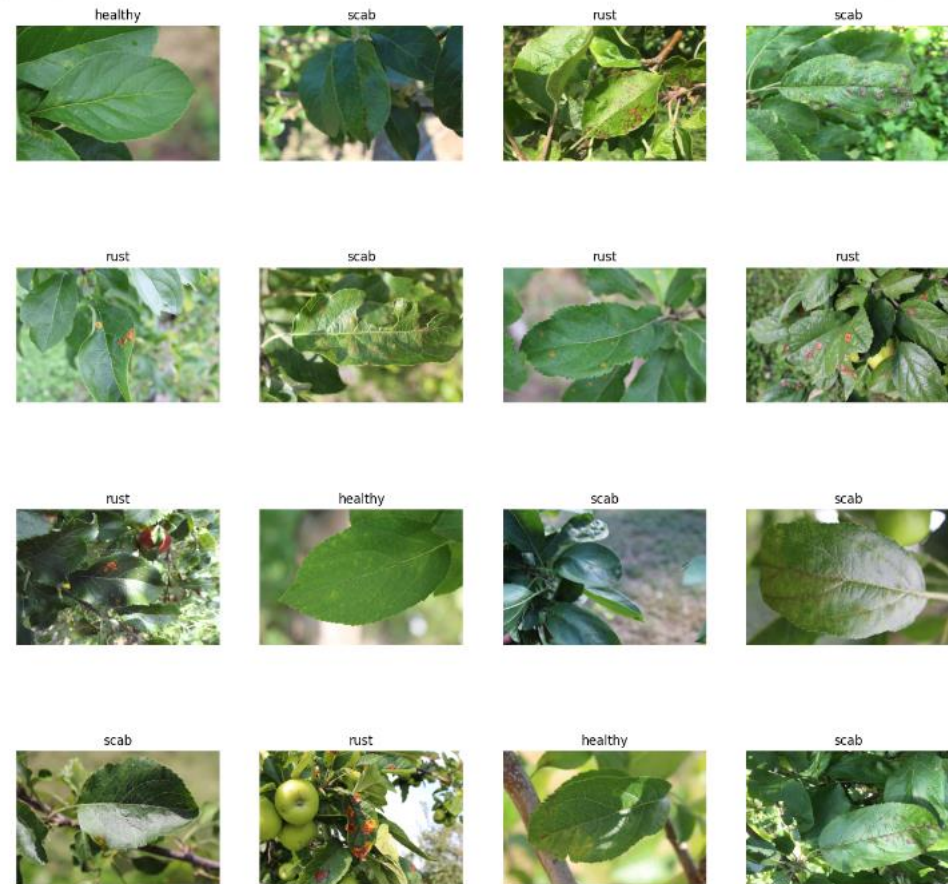
(Predictions)

```
In [ ]: nrows = 4
ncols = 4

fig = plt.gcf()
fig.set_size_inches(ncols*4, nrows*4)

for i, (idx, row) in enumerate(df_out.sample(nrows*ncols).iterrows()):
    filepath = os.path.join(SOURCE, idx + ".jpg")
    sp = plt.subplot(nrows, ncols, i + 1)
    sp.axis('Off') # Don't show axes (or gridlines)
    img = mpimg.imread(filepath)
    plt.title(df_out.columns[np.argmax(row)])
    plt.imshow(img)

plt.show()
```



A DenseNet is a type of convolutional neural network that utilises dense connections between layers, through Dense Blocks, where we connect all layers (with matching feature-map sizes) directly with each other.

3- Dense Net
Pretrained Model

3- Dense Net

Pretrained Model

(Import dense model -Transfer
learning – compile - fit)

```
from tensorflow.keras.applications import DenseNet201
def dense_net_model(trainable_weights=False, weights_path=None):

    tf.keras.backend.clear_session()

    dense_net = DenseNet201(input_shape=(256, 256, 3), weights="imagenet", include_top=False)

    for layer in dense_net.layers:
        layer.trainable=trainable_weights

    model = tf.keras.models.Sequential([dense_net,
                                        tf.keras.layers.GlobalAveragePooling2D(),
                                        tf.keras.layers.Dense(128, activation='relu'),
                                        tf.keras.layers.Dropout(0.3),
                                        tf.keras.layers.Dense(4, activation='softmax')])

    if weights_path:
        model.load_weights(weights_path)

    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate_scheduler)
    model.compile(loss="categorical_crossentropy", optimizer=optimizer, metrics=['accuracy'])

    return model

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=3)
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint("dense.h5", save_best_only=True)

dense_net_transfer = dense_net_model(trainable_weights=True)

dense_net_transfer_history = dense_net_transfer.fit(train_data_multi, validation_data=val_data_multi, epochs=25, steps_per_ep
```

Dense Net

Pretrained Model

(Import dense model -Transfer learning – compile - fit)

```
def dense_net_model(trainable_weights=False, weights_path=None):  
    tf.keras.backend.clear_session()  
  
    dense_net = DenseNet201(input_shape=(256, 256, 3), weights="imagenet", include_top=False)  
  
    for layer in dense_net.layers:  
        layer.trainable=trainable_weights  
  
    model = tf.keras.models.Sequential([dense_net,  
                                        tf.keras.layers.GlobalAveragePooling2D(),  
                                        tf.keras.layers.Dense(128, activation='relu'),  
                                        tf.keras.layers.Dropout(0.3),  
                                        tf.keras.layers.Dense(4, activation='softmax')])  
  
    if weights_path:  
        model.load_weights(weights_path)  
  
    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate_scheduler)  
    model.compile(loss="categorical_crossentropy", optimizer=optimizer, metrics=['accuracy'])  
  
    return model  
  
early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=3)  
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint("dense.h5", save_best_only=True)  
  
dense_net_transfer = dense_net_model(trainable_weights=True)  
  
dense_net_transfer_history = dense_net_transfer.fit(train_data_multi, validation_data=val_data_multi, epochs=25, steps_per_ep
```

```
Epoch 1/25  
32/32 [-----] - 218s 2s/step - loss: 0.6594 - accuracy: 0.7178 - val_loss: 1.6250 - val_accuracy: 0.4824  
Epoch 2/25  
32/32 [-----] - 57s 2s/step - loss: 0.2936 - accuracy: 0.8959 - val_loss: 0.4921 - val_accuracy: 0.8338  
Epoch 3/25  
32/32 [-----] - 51s 2s/step - loss: 0.2444 - accuracy: 0.9839 - val_loss: 0.3598 - val_accuracy: 0.8692  
Epoch 4/25  
32/32 [-----] - 54s 2s/step - loss: 0.1365 - accuracy: 0.9521 - val_loss: 0.3284 - val_accuracy: 0.8893  
Epoch 5/25  
32/32 [-----] - 52s 2s/step - loss: 0.0981 - accuracy: 0.9697 - val_loss: 0.2828 - val_accuracy: 0.9395  
Epoch 6/25  
32/32 [-----] - 53s 2s/step - loss: 0.0956 - accuracy: 0.9717 - val_loss: 0.1788 - val_accuracy: 0.9497  
Epoch 7/25  
32/32 [-----] - 53s 2s/step - loss: 0.0914 - accuracy: 0.9710 - val_loss: 0.1631 - val_accuracy: 0.9557  
Epoch 8/25  
32/32 [-----] - 55s 2s/step - loss: 0.0840 - accuracy: 0.9688 - val_loss: 0.1681 - val_accuracy: 0.9537  
Epoch 9/25  
32/32 [-----] - 58s 2s/step - loss: 0.0640 - accuracy: 0.9820 - val_loss: 0.1597 - val_accuracy: 0.9437  
Epoch 10/25  
32/32 [-----] - 51s 2s/step - loss: 0.0737 - accuracy: 0.9780 - val_loss: 0.1388 - val_accuracy: 0.9577  
Epoch 11/25  
32/32 [-----] - 58s 2s/step - loss: 0.0733 - accuracy: 0.9820 - val_loss: 0.1218 - val_accuracy: 0.9598  
Epoch 12/25  
32/32 [-----] - 51s 2s/step - loss: 0.0738 - accuracy: 0.9736 - val_loss: 0.1112 - val_accuracy: 0.9658  
Epoch 13/25  
32/32 [-----] - 49s 2s/step - loss: 0.0666 - accuracy: 0.9824 - val_loss: 0.1234 - val_accuracy: 0.9678  
Epoch 14/25  
32/32 [-----] - 58s 2s/step - loss: 0.0828 - accuracy: 0.9780 - val_loss: 0.1070 - val_accuracy: 0.9718  
Epoch 15/25  
32/32 [-----] - 52s 2s/step - loss: 0.0690 - accuracy: 0.9830 - val_loss: 0.1022 - val_accuracy: 0.9718  
Epoch 16/25  
32/32 [-----] - 49s 2s/step - loss: 0.0751 - accuracy: 0.9785 - val_loss: 0.1027 - val_accuracy: 0.9718  
Epoch 17/25  
32/32 [-----] - 51s 2s/step - loss: 0.0770 - accuracy: 0.9770 - val_loss: 0.0930 - val_accuracy: 0.9718  
Epoch 18/25  
32/32 [-----] - 53s 2s/step - loss: 0.0564 - accuracy: 0.9873 - val_loss: 0.0902 - val_accuracy: 0.9738  
Epoch 19/25  
32/32 [-----] - 48s 2s/step - loss: 0.0540 - accuracy: 0.9883 - val_loss: 0.0976 - val_accuracy: 0.9738  
Epoch 20/25  
32/32 [-----] - 49s 2s/step - loss: 0.0748 - accuracy: 0.9790 - val_loss: 0.1067 - val_accuracy: 0.9698  
Epoch 21/25  
32/32 [-----] - 48s 2s/step - loss: 0.0780 - accuracy: 0.9790 - val_loss: 0.0974 - val_accuracy: 0.9598
```

Dense Net


Pretrained Model

(train and validation accuracy
visualization – load model -
evaluation)

```
In [ ]: import matplotlib.pyplot as plt
acc = dense_net_transfer_history.history['accuracy']
val_acc = dense_net_transfer_history.history['val_accuracy']
loss = dense_net_transfer_history.history['loss']
val_loss = dense_net_transfer_history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend(loc=0)
```



```
In [ ]: #model = tf.keras.models.load_model("dense.h5") # rollback to best model
dense_net_transfer.evaluate(val_data_multi)
```

16/16 [=====] - 14s 837ms/step - loss: 0.0915 - accuracy: 0.9678

Out[18]: [0.09147188812494278, 0.9678068161010742]

Dense Net

Pretrained Model

(read test data – make predictions –
Save submission file)

```
In [ ]: test_df = pd.read_csv("/kaggle/input/plant-pathology-2020-fgvc7/test.csv")
test_df['img_name'] = test_df['image_id'] + ".jpg"

test_datagen = ImageDataGenerator(rescale=1/255)

test_generator = test_datagen.flow_from_dataframe(dataframe=test_df,
                                                  directory="/kaggle/input/plant-pathology-2020-fgvc7/images/",
                                                  x_col="img_name",
                                                  y_col=None,
                                                  target_size=(256, 256),
                                                  class_mode=None,
                                                  batch_size=3,
                                                  shuffle=False,
                                                  seed=42)

test_generator.reset()
```

Found 1821 validated image filenames.

```
In [ ]: preds = dense_net_transfer.predict_generator(test_generator, verbose=1, steps=607)

preds_df = pd.DataFrame(preds, columns=["healthy", "multiple_diseases", "rust", "scab"])

submission = pd.concat([test_df.image_id, preds_df], axis=1)
```

607/607 [=====] - 46s 67ms/step

Out[20]:

	image_id	healthy	multiple_diseases	rust	scab
0	Test_0	0.000111	0.002933	0.996945	0.000011
1	Test_1	0.000018	0.001202	0.998775	0.000004
2	Test_2	0.177084	0.414771	0.087757	0.320388
3	Test_3	0.999806	0.000059	0.000040	0.000095
4	Test_4	0.000008	0.000232	0.999758	0.000002

```
In [ ]: submission.to_csv("submission_DenseNet.csv", index=False)
```

Dense Net

Pretrained Model

(read test data – make predictions –
Save submission file)

```
In [ ]: test_df = pd.read_csv("/kaggle/input/plant-pathology-2020-fgvc7/test.csv")
test_df['img_name'] = test_df['image_id'] + ".jpg"

test_datagen = ImageDataGenerator(rescale=1/255)

test_generator = test_datagen.flow_from_dataframe(dataframe=test_df,
                                                  directory="/kaggle/input/plant-pathology-2020-fgvc7/images/",
                                                  x_col="img_name",
                                                  y_col=None,
                                                  target_size=(256, 256),
                                                  class_mode=None,
                                                  batch_size=3,
                                                  shuffle=False,
                                                  seed=42)

test_generator.reset()
```

Found 1821 validated image filenames.

```
In [ ]: preds_df = pd.DataFrame(preds, columns=["healthy", "multiple_diseases", "rust", "scab"])

submission = pd.concat([test_df.image_id, preds_df], axis=1)

submission.head()
```

607/607 [=====] - 46s 67ms/step

Out[20]:

	image_id	healthy	multiple_diseases	rust	scab
0	Test_0	0.000111	0.002933	0.996945	0.000011
1	Test_1	0.000018	0.001202	0.998775	0.000004
2	Test_2	0.177084	0.414771	0.087757	0.320388
3	Test_3	0.999806	0.000059	0.000040	0.000095
4	Test_4	0.000008	0.000232	0.999758	0.000002

```
In [ ]: submission.to_csv("submission_DenseNet.csv", index=False)
```

```
In [ ]: den_model = tf.keras.models.load_model("dense.h5") # rollback to best model
den_model.evaluate(val_data_multi)
```

16/16 [=====] - 17s 826ms/step - loss: 0.0863 - accuracy: 0.9799

Out[23]: [0.08627106249332428, 0.9798792600631714]

Made by

Karima Sobhi

Haidy Ashraf

Nada Abdallah

Thank You

